

Nervous-System-Aware Platform-Risk Intelligence

Nervous-System-Aware Platform-Risk Intelligence

Version 2026.08.0 · evidence cutoff 2026-08-03 · frozen pre-release

A dated monitoring and response aid, not proof of suppression, diagnosis of any person, legal or medical advice, platform access, appeal-success assurance, or restored-distribution guarantee. Platform ranking and enforcement systems are opaque and change; reverify every policy, price, account surface, and decision before consequential use.

Operating proposition

Platform dependence creates uncertainty, loss of control, and social-threat interpretation. The answer is not to diagnose a creator or anthropomorphize a ranking system. It is to design a calm, evidence-preserving workflow: separate observation from inference, establish comparable baselines, inspect official status surfaces, test alternatives, use independent review, and escalate proportionally.

Statistical method

For at least five comparable historical periods, calculate the arithmetic mean and sample standard deviation. The illustrative lower control limit is $\max(0, \text{mean} - 3s)$. Falling below it is a screening signal, not causal proof. Do not pool unlike formats, audiences, paid campaigns, or lifecycle stages. Recalculate only under a documented baseline-change rule; never move the baseline merely to erase an anomaly.

Decision language

Say “observed decline,” “below the declared control limit,” “official eligibility record,” or “hypothesis under test.” Do not say shadowbanned, suppressed, penalized, or diagnosed unless the exact claim is supported by a preserved official record or clearly labeled as an allegation.

Exact 45-minute monthly review

0:00–05:00 — regulate the room: restate the no-causal-certainty boundary; assign facilitator and recorder. 05:00–15:00 — verify data health, exports, dates, content IDs, campaign effects, outages, and missingness. 15:00–25:00 — recalculate baselines and control limits from comparable content; flag signals without naming a cause. 25:00–33:00 — inspect native account-status, eligibility, enforcement, and appeal surfaces; preserve dated evidence. 33:00–39:00 — test alternative explanations: content mix, seasonality, paid activity, audience, creative, frequency, tracking, and ordinary ranking variation. 39:00–43:00 — classify Observe / Investigate / Respond; assign owner, evidence gap, next test, and stop condition. 43:00–45:00 — schedule follow-up, export the record, and confirm no diagnosis, hidden-system certainty, appeal success, or restored reach was promised.

References

See `source-ledger.json` and `policy-ledger.csv`. Official platform sources are point-in-time records and must be reverified. NIST AI RMF is voluntary and its functions are adapted here without endorsement.