

Neuroscience-of-Trust Content Architecture

Neuroscience-of-Trust Content Architecture

Version 2026.08.0 · source-review cutoff 2026-08-03 · frozen pre-release

What this package does

EDSA helps a human editor make two appraisals before publishing: what an audience needs to evaluate a message without coercion, and what a platform may need to classify and distribute it. Those appraisals are related but not interchangeable. The method converts them into four observable editorial moves: Evidence, Discussion, Self-Determination, and Action.

EDSA is an editorial decision architecture. It cannot diagnose a nervous system, predict an individual's response, establish scientific causation, guarantee trust, guarantee platform distribution, or replace legal, medical, scientific, accessibility, or platform-policy review.

Dual appraisal

Audience appraisal asks whether the opening creates intelligibility or avoidable threat; whether evidence is inspectable; whether uncertainty is proportionate; whether the audience retains choice; and whether the requested action is reversible and bounded. Platform appraisal asks whether the content is eligible, understandable outside the creator's network, internally coherent across text and media, relevant to likely viewers, and free of policy or quality signals that may limit recommendation. Platform documentation describes inputs and eligibility rules, not a reliable recipe for reach.

EDSA

1. **Evidence** — name the proposition, show the best available authority beside it, preserve date, scope, and uncertainty, and distinguish observation from inference.
2. **Discussion** — explain competing interpretations, mechanism limits, context, and what would change the conclusion. Do not manufacture balance where evidence is one-sided.

3. **Self-Determination** — return choice to the audience: identify who the content is for, offer a pause or exit, disclose material interests, and avoid shame, urgency, or identity pressure.
4. **Action** — request one proportionate, specific, reversible next step. State what the action cannot establish and who must review consequential decisions.

Five hard-never rules

1. Never label an emotional response as a diagnosis or claim to regulate a viewer's nervous system.
2. Never convert a neuroscience association into a causal promise about trust, attention, memory, persuasion, or behavior.
3. Never imply that a format, hook, retention tactic, or policy-compliant post is guaranteed distribution.
4. Never hide uncertainty, sponsorship, material interest, eligibility limits, or a consequential professional boundary.
5. Never use fear, shame, manufactured urgency, identity threat, or false scarcity to remove meaningful audience choice.

Operating sequence

Start with the content brief and claim map. Complete the dual-appraisal crosswalk. Draft the twelve-minute spine or a shorter proportionate variant. Run all four EDSA passes. Check the five hard-never rules. Verify every platform statement against its dated source record. Ask a human reviewer to record approve, revise, reject, or escalate. Archive the exact version used.

Evaluation, not prediction

Measure observable editorial properties: source resolution, claim fidelity, uncertainty, choice-preserving language, context, accessibility, policy status, and a bounded CTA. Post-publication analytics may describe what happened for that content and audience. They do not prove why it happened, predict an individual response, or validate a neuroscience mechanism.

When to hire help

Use qualified scientific review for mechanism claims; legal or regulatory review for regulated or consequential statements; accessibility expertise before claiming conformance; clinical review before discussing individual symptoms or treatment; and platform counsel or current first-party policy documentation when eligibility matters. The package is an editorial aid, not professional advice or assurance.